AI Accelerator Program

Analytics Vidhya

5 RAG Examples from Real Companies

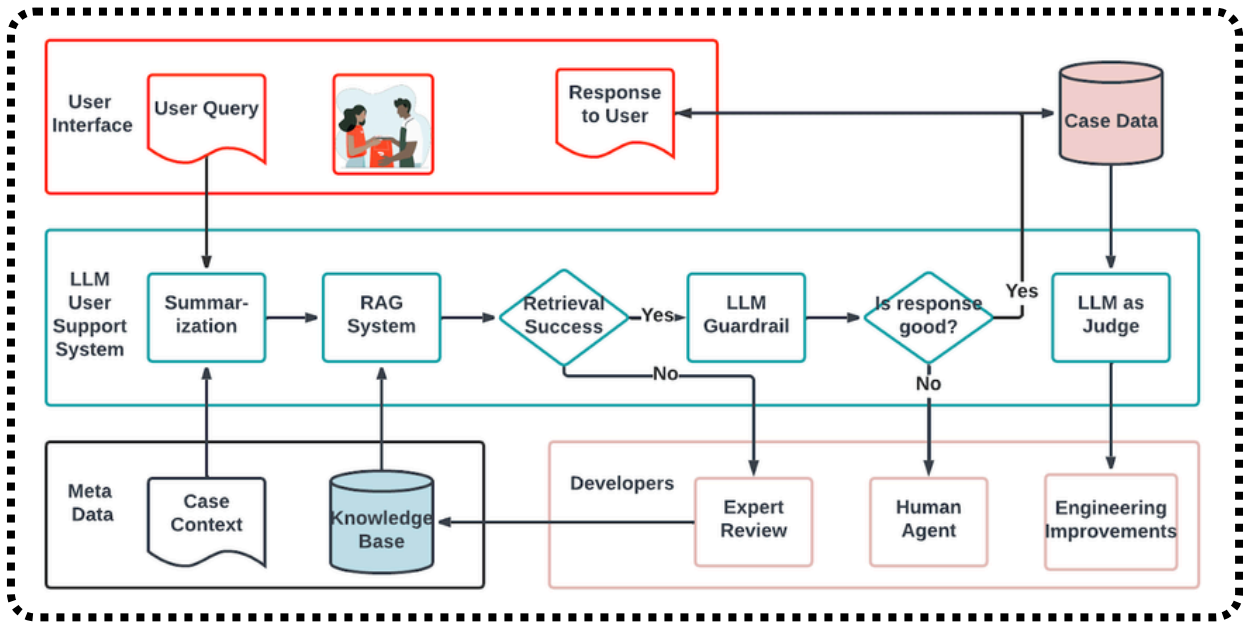

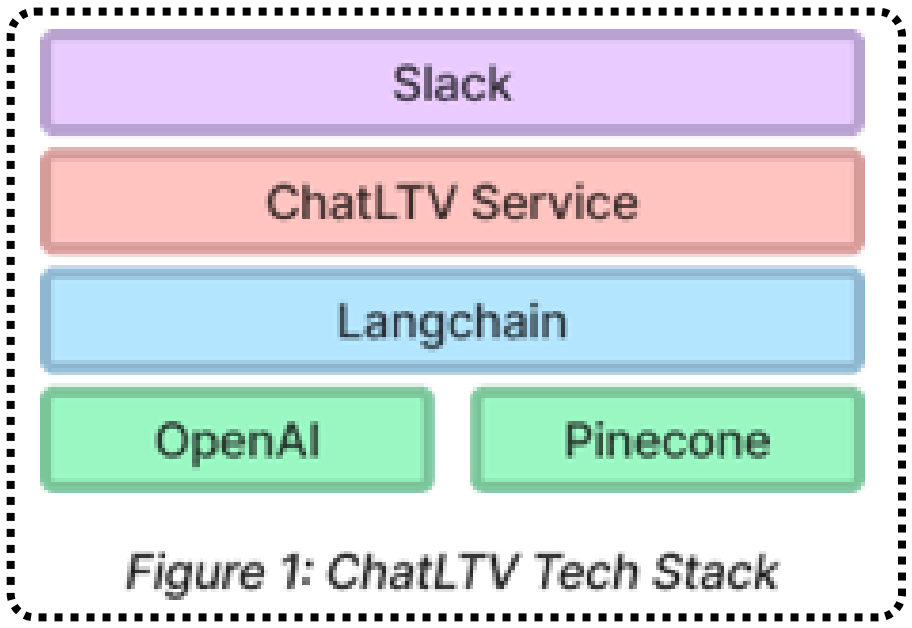

Figure 1: ChatLTV Tech Stack

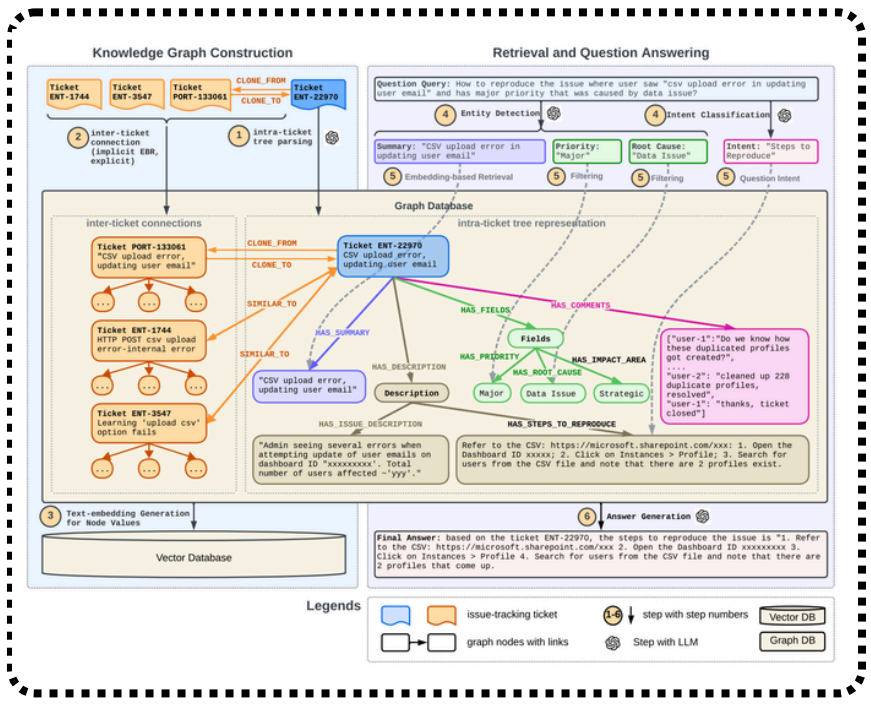

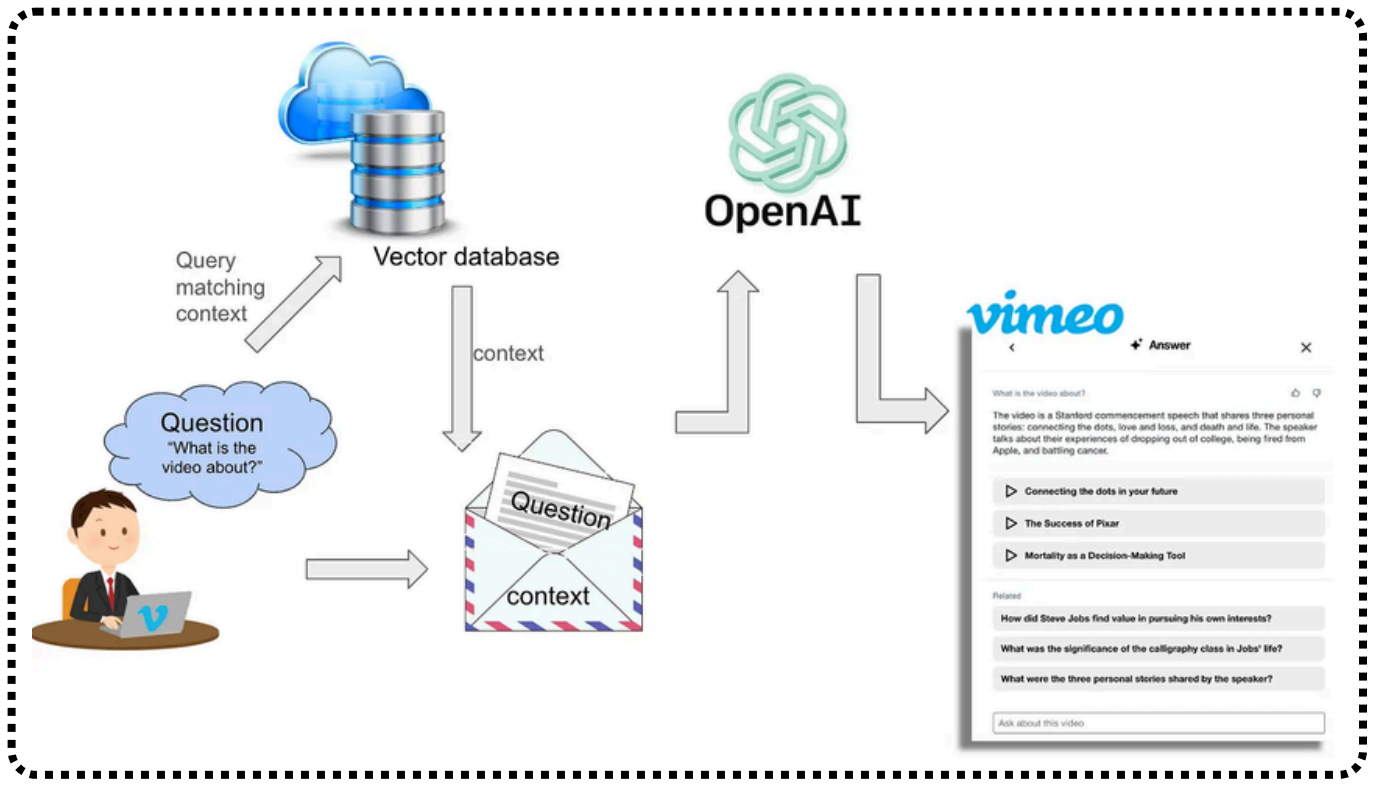

DOORDASH

LinkedIn

Bell

vimeo

Harvard Business School

CHAT LTV

Pinterest

DoorDash

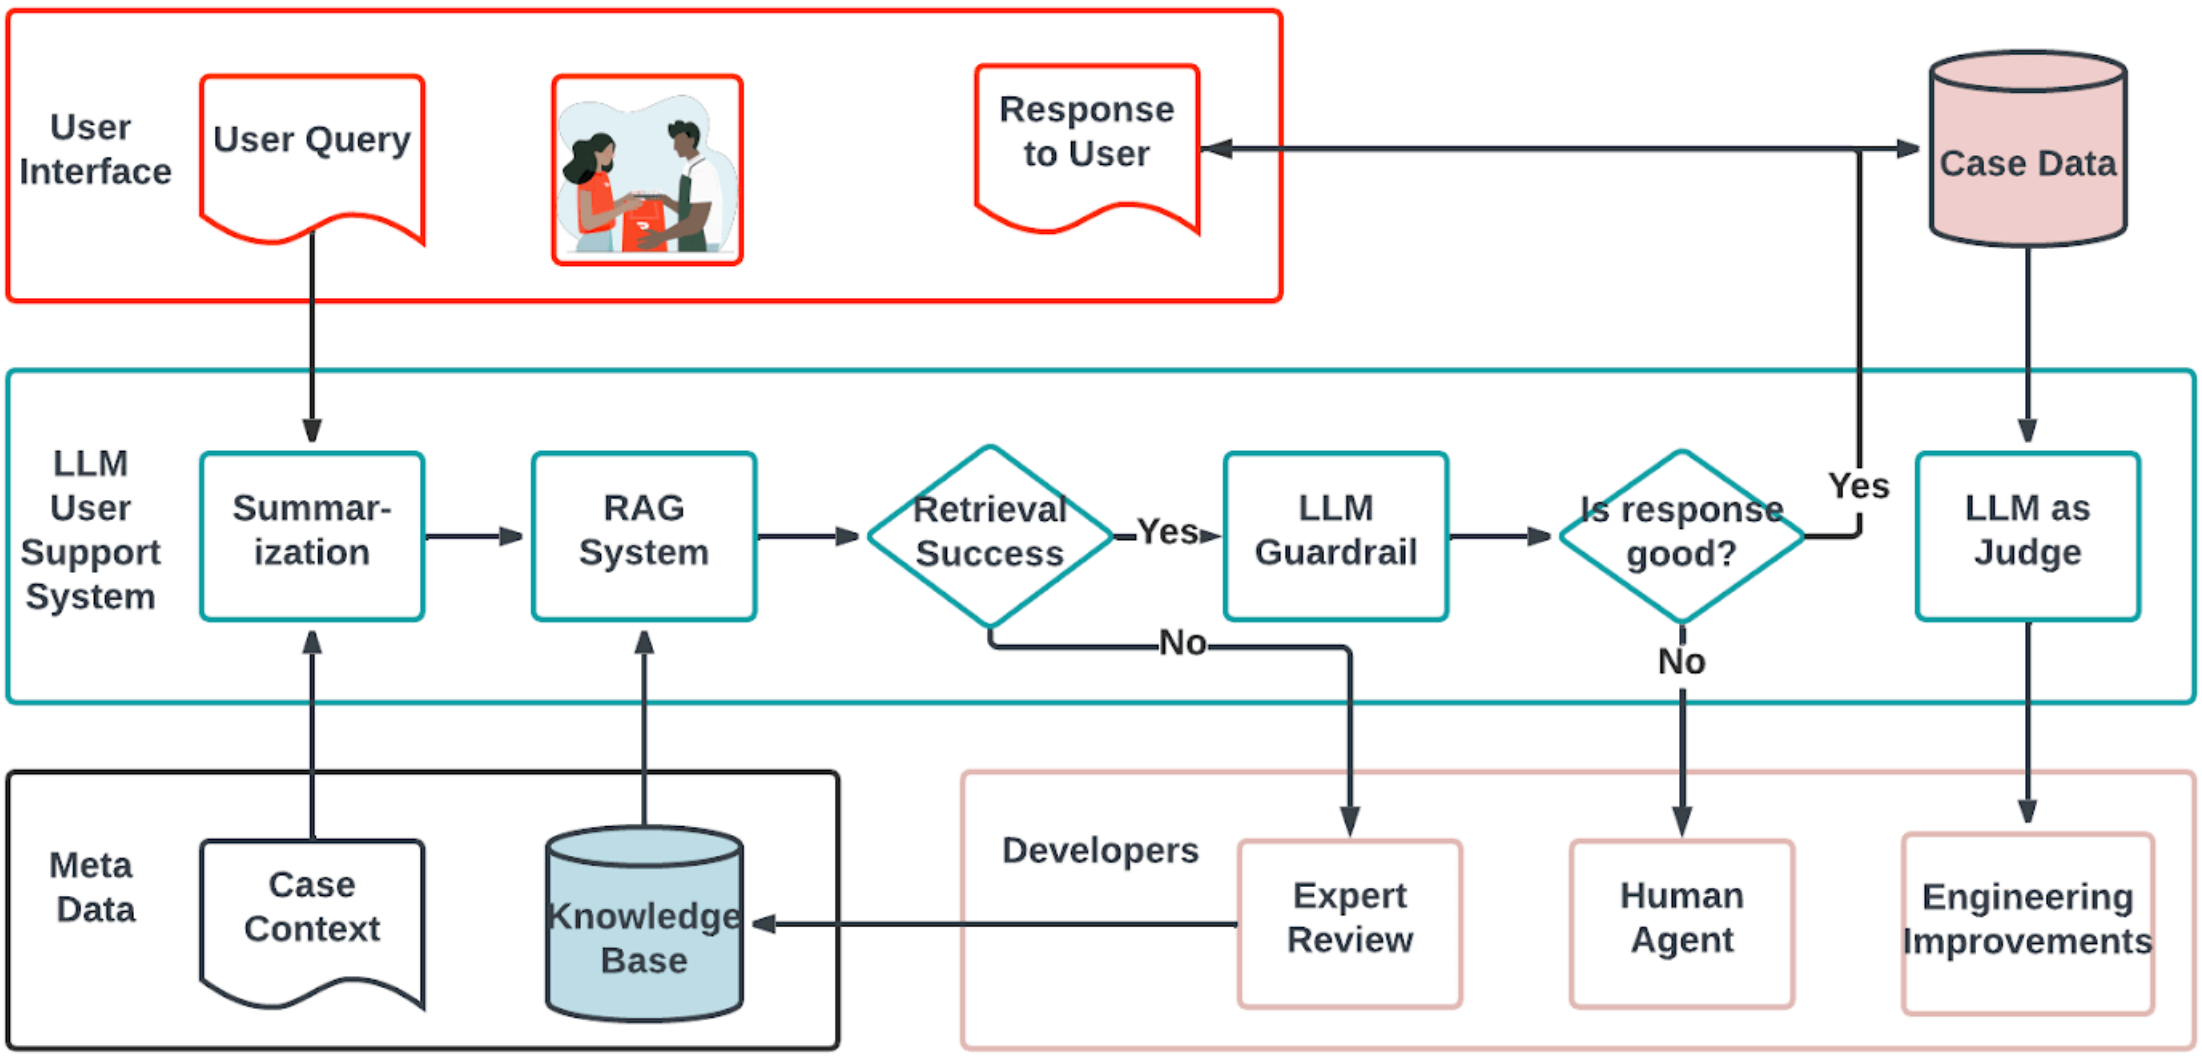

DoorDash improves delivery support using a RAG-based chatbot built with three main components: a **RAG system**, an **LLM Guardrail**, and an **LLM Judge**.

When a Dasher reports an issue, the system first summarizes the conversation to clearly identify the core problem. Based on this summary, the RAG system retrieves the most relevant knowledge base articles and previously resolved cases. This retrieved context is then passed to an LLM, which generates a coherent and context-aware response tailored to the Dasher's query.

To ensure response quality, DoorDash implemented an LLM Guardrail system. This acts as an online monitoring layer that evaluates each LLM-generated response for factual accuracy and policy compliance. It helps reduce hallucinations and filters out outputs that violate company guidelines.

To track long-term system performance, DoorDash uses an LLM Judge. The LLM Judge evaluates the chatbot across five metrics: retrieval correctness, response accuracy, grammar and language quality, contextual coherence, and relevance to the Dasher's request.

Customer Tech Support

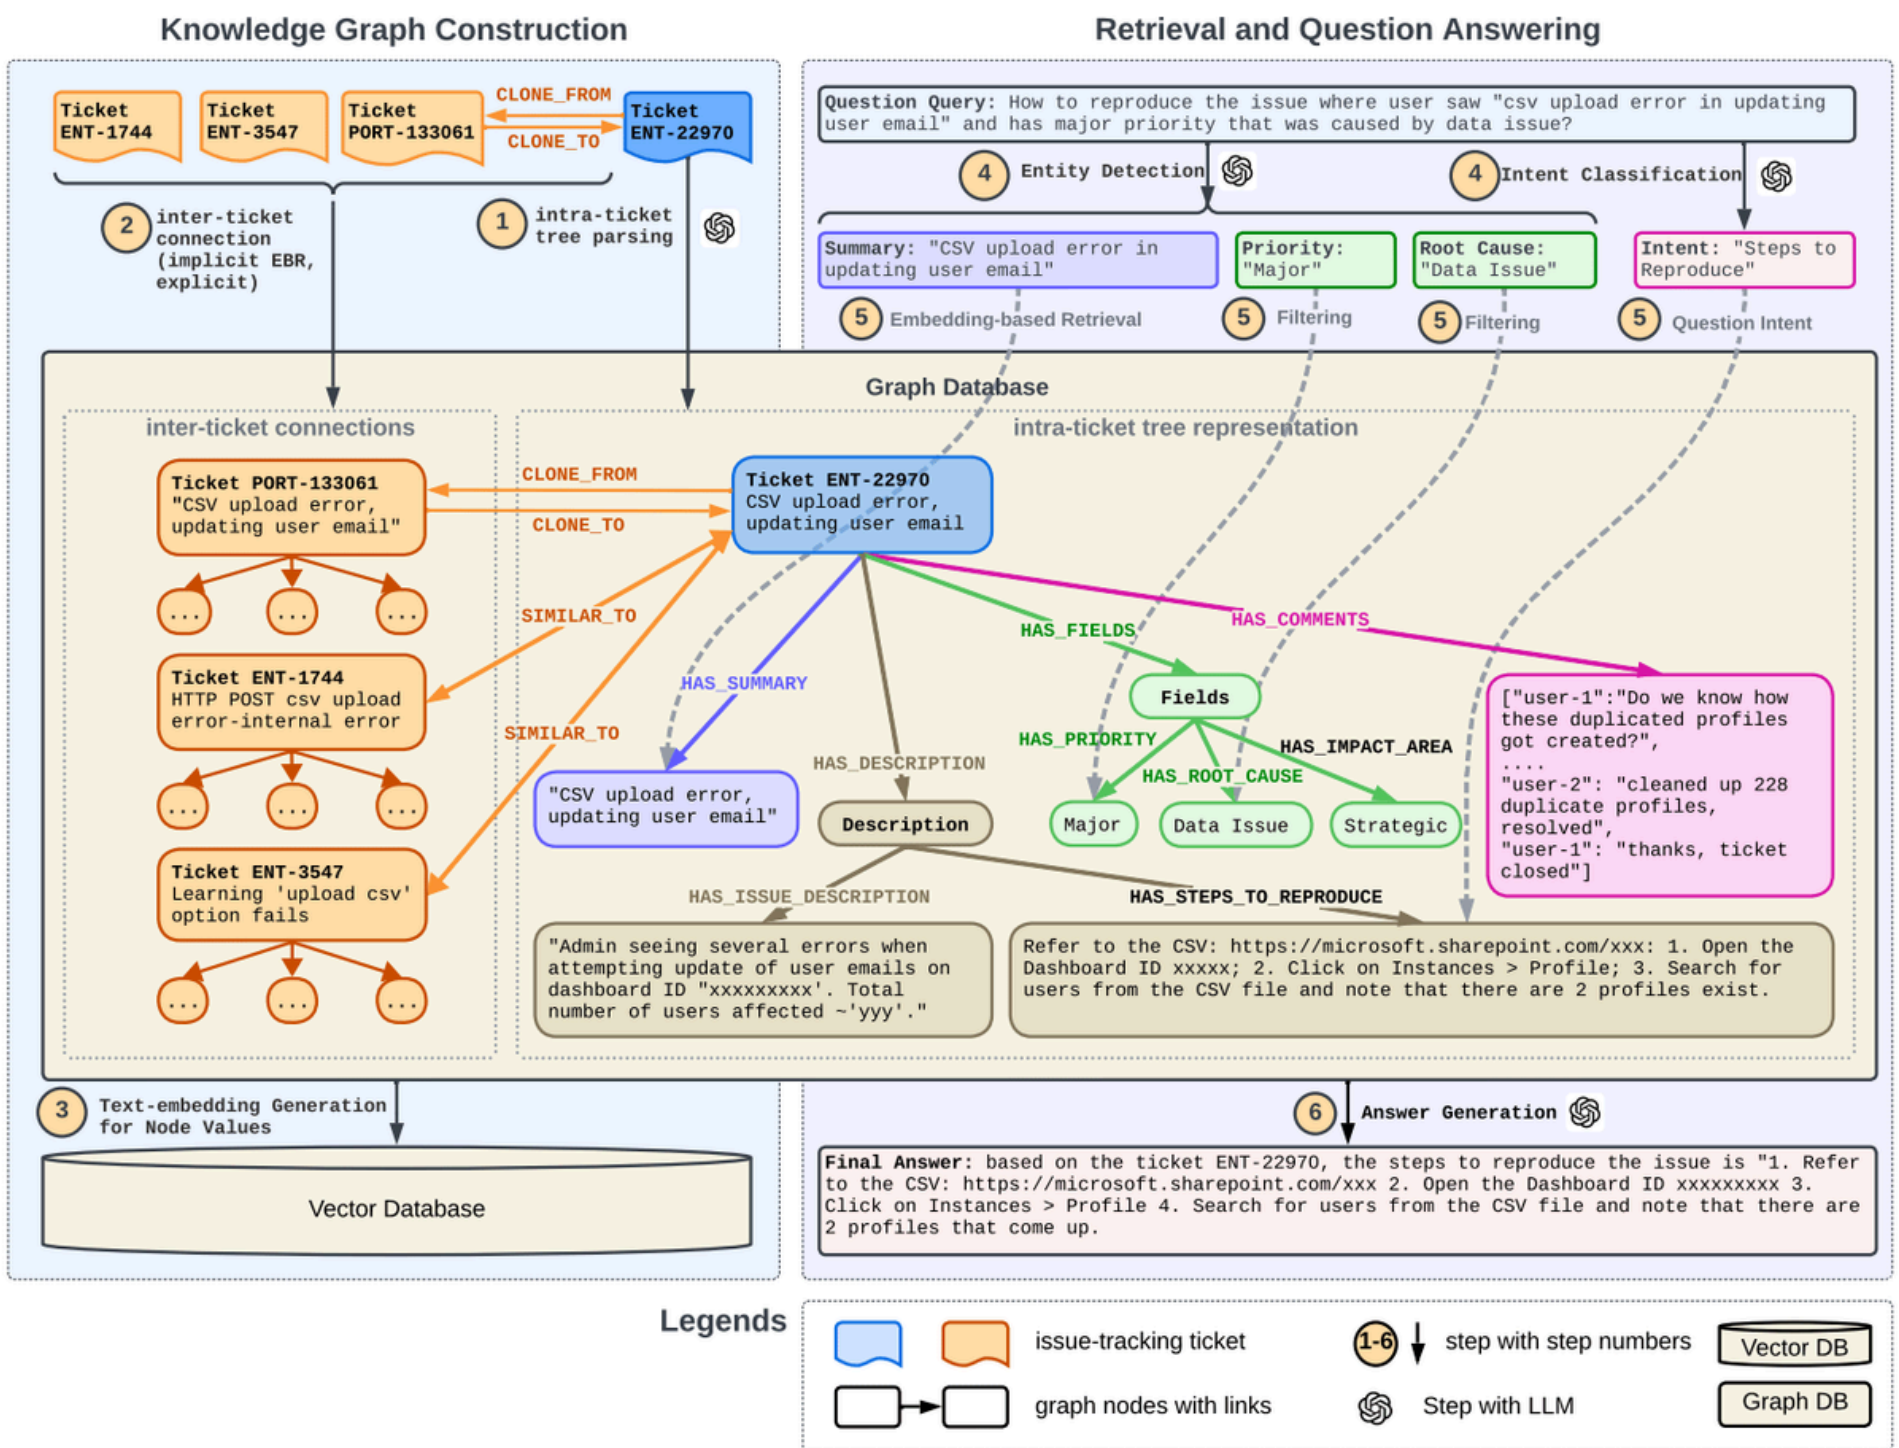

LinkedIn introduced a customer service question-answering system that combines RAG with a knowledge graph.

Instead of treating past issue-tracking tickets as plain text, the system builds a knowledge graph from historical issues. This graph captures both the structure within each issue and the relationships between different issues.
When a user submits a question, the system parses the query and retrieves relevant sub-graphs from the knowledge graph. These sub-graphs are then used as context for the RAG system to generate an answer. By relying on structured graph data instead of simple text chunks, the approach reduces problems caused by text segmentation and improves retrieval accuracy.

This system has been deployed within LinkedIn’s customer service team and has reduced the median issue resolution time by 28.6%.

Internal policies chatbot

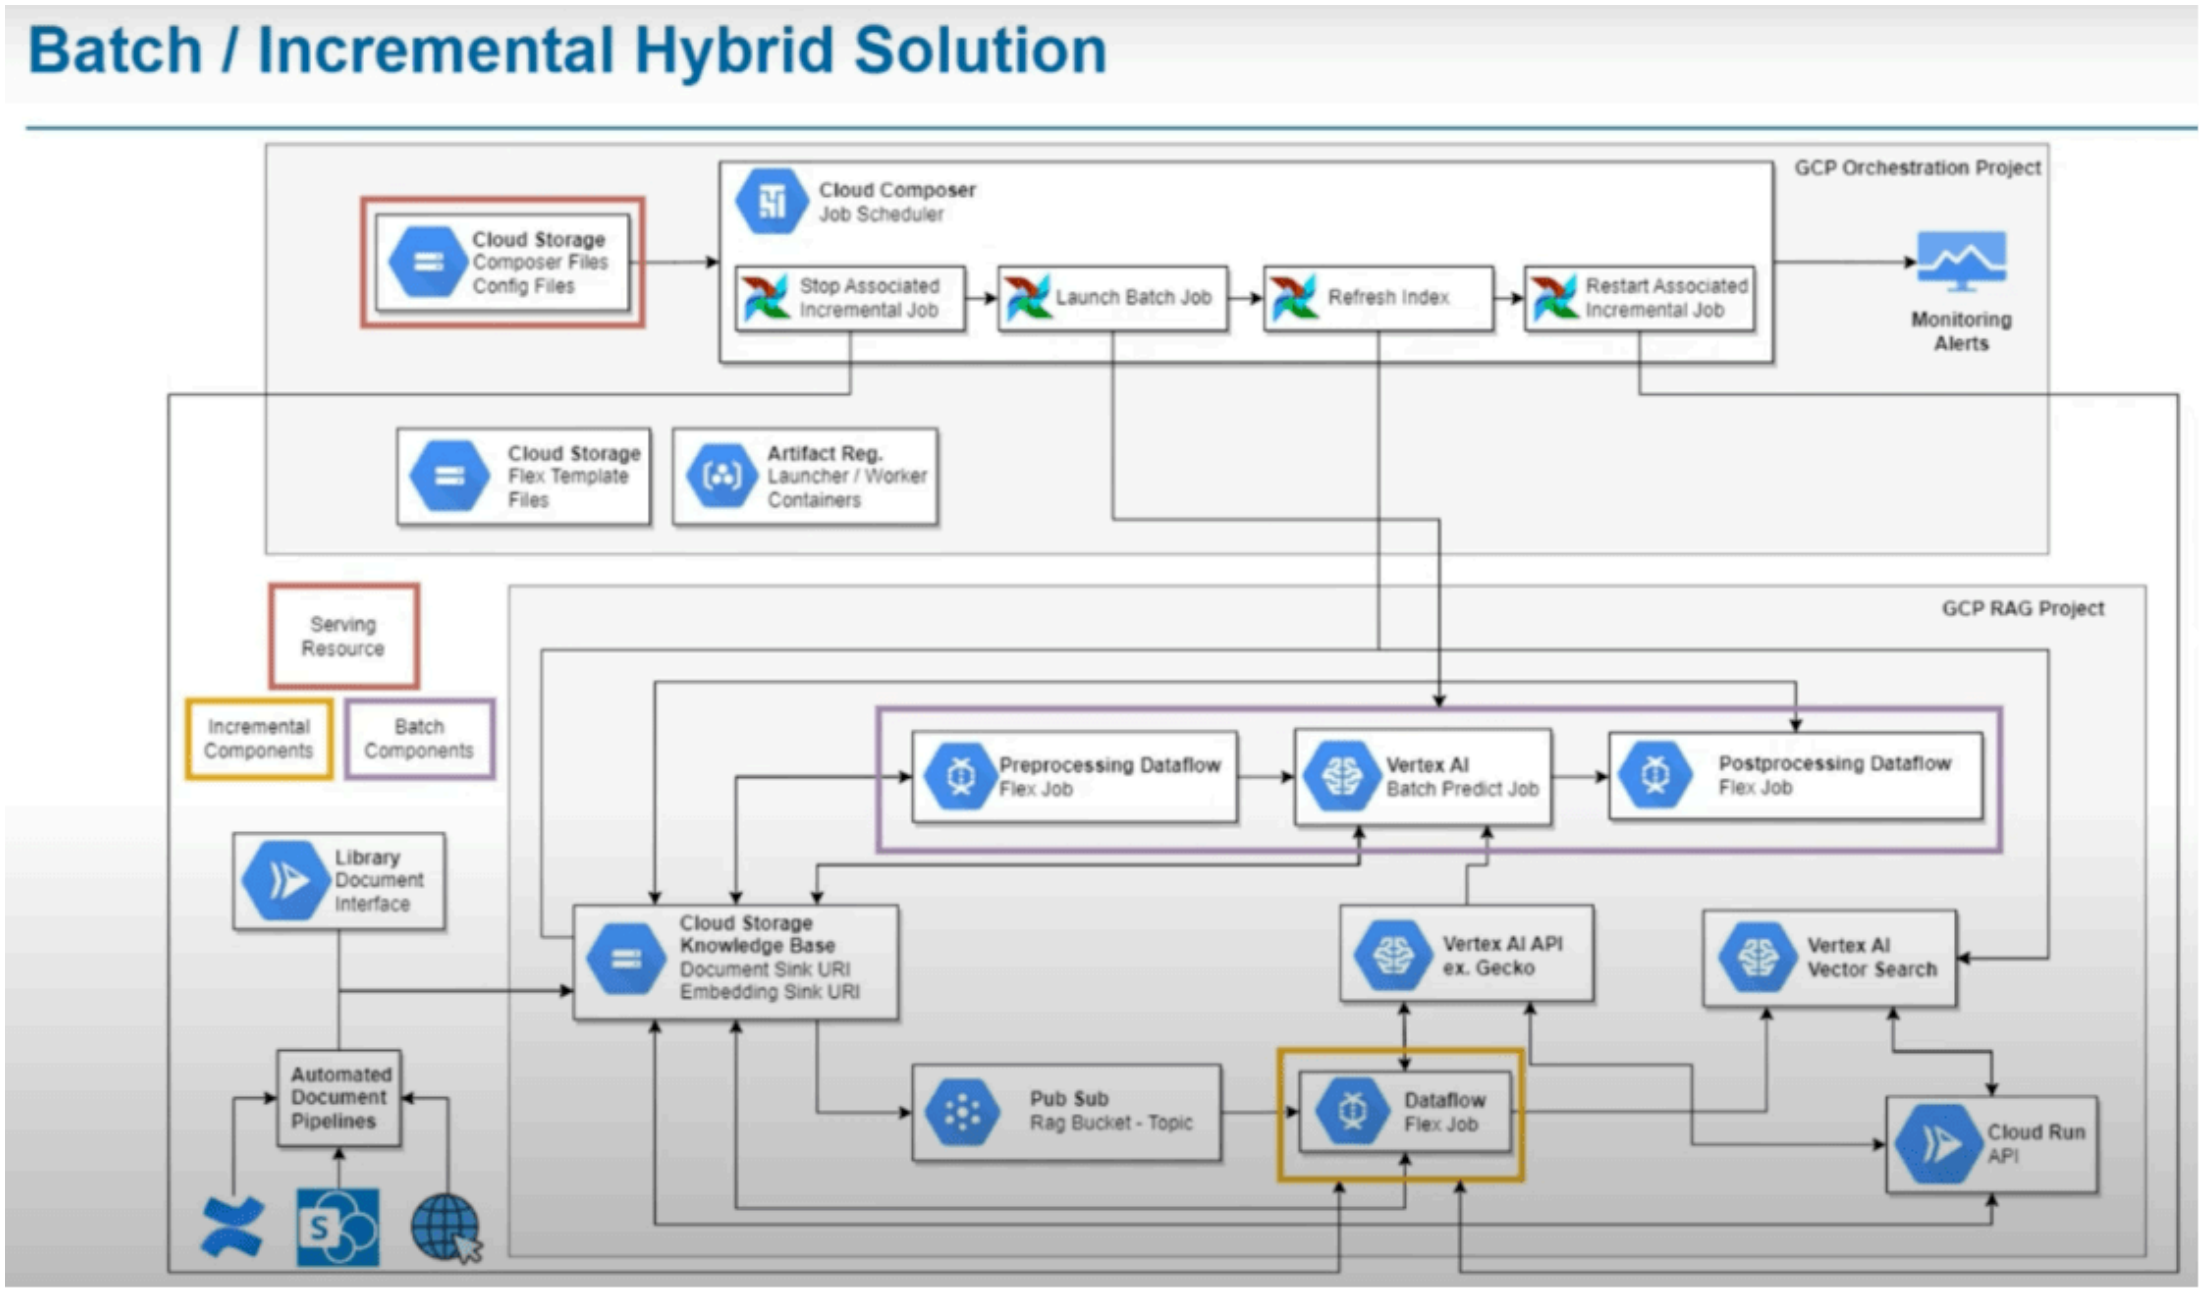

Bell, a telecommunications company, uses RAG to improve its knowledge management system and make sure employees can easily access the latest company policies.

As part of this setup, Bell built a dedicated knowledge management component within its RAG system. It uses modular document embedding pipelines to process and index raw documents from multiple sources efficiently.

The system supports both batch updates and incremental updates, so new documents can be added without rebuilding everything from scratch. It also automatically refreshes indexes whenever documents are added, updated, or removed from their original source.

Each component of the system is designed as an independent service. Bell follows DevOps practices to develop, deploy, and maintain the entire RAG-based solution.

AI Professor

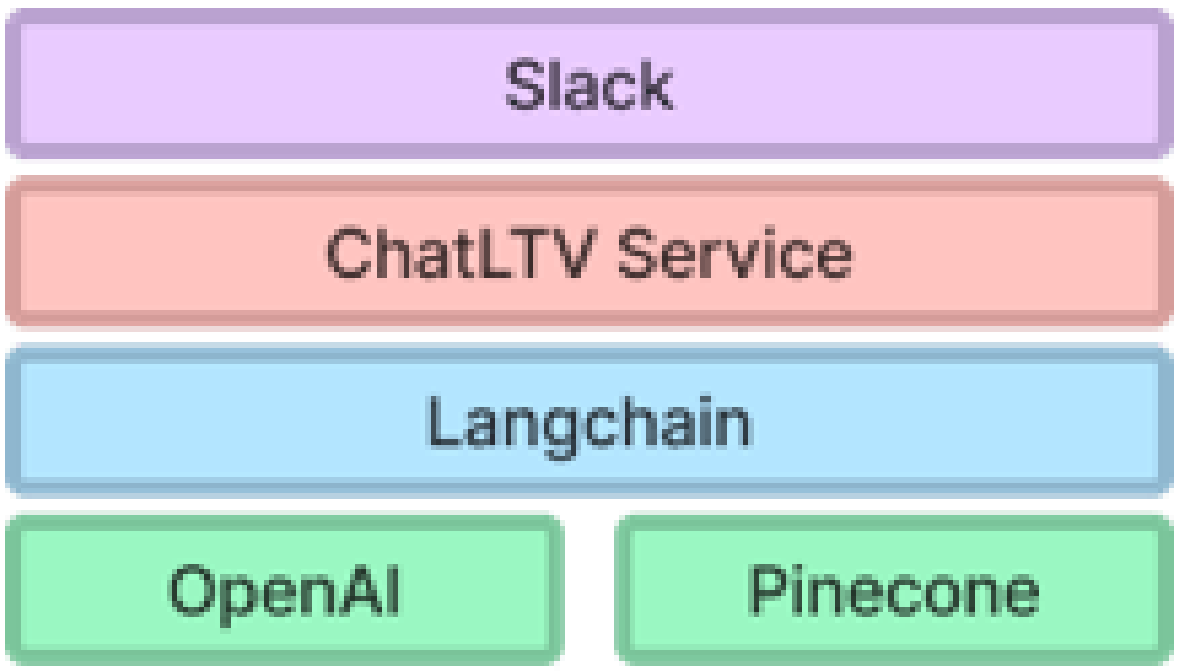

Figure 1: ChatLTV Tech Stack

At Harvard Business School, Senior Lecturer Jeffrey Bussgang built a RAG-based AI faculty chatbot called ChatLTV to support his entrepreneurship course.

ChatLTV helps students prepare for class by explaining complex concepts, finding additional insights on case studies, and answering administrative questions.

The chatbot was trained on the course corpus, including case studies, teaching notes, books, blog posts, and past Q&A from the course's Slack channel.

It is integrated directly into Slack, so students can interact with it in both private and public channels.

When a student asks a question, the system retrieves relevant context from a vector database. The most relevant content chunks are then sent, along with the query, to an LLM through the OpenAI API. The LLM uses this retrieved context to generate a grounded response.

To maintain response quality, the course team used both manual and automated evaluation. They implemented an LLM Judge to compare chatbot outputs against ground-truth answers and assign a quality score.

Video summaries

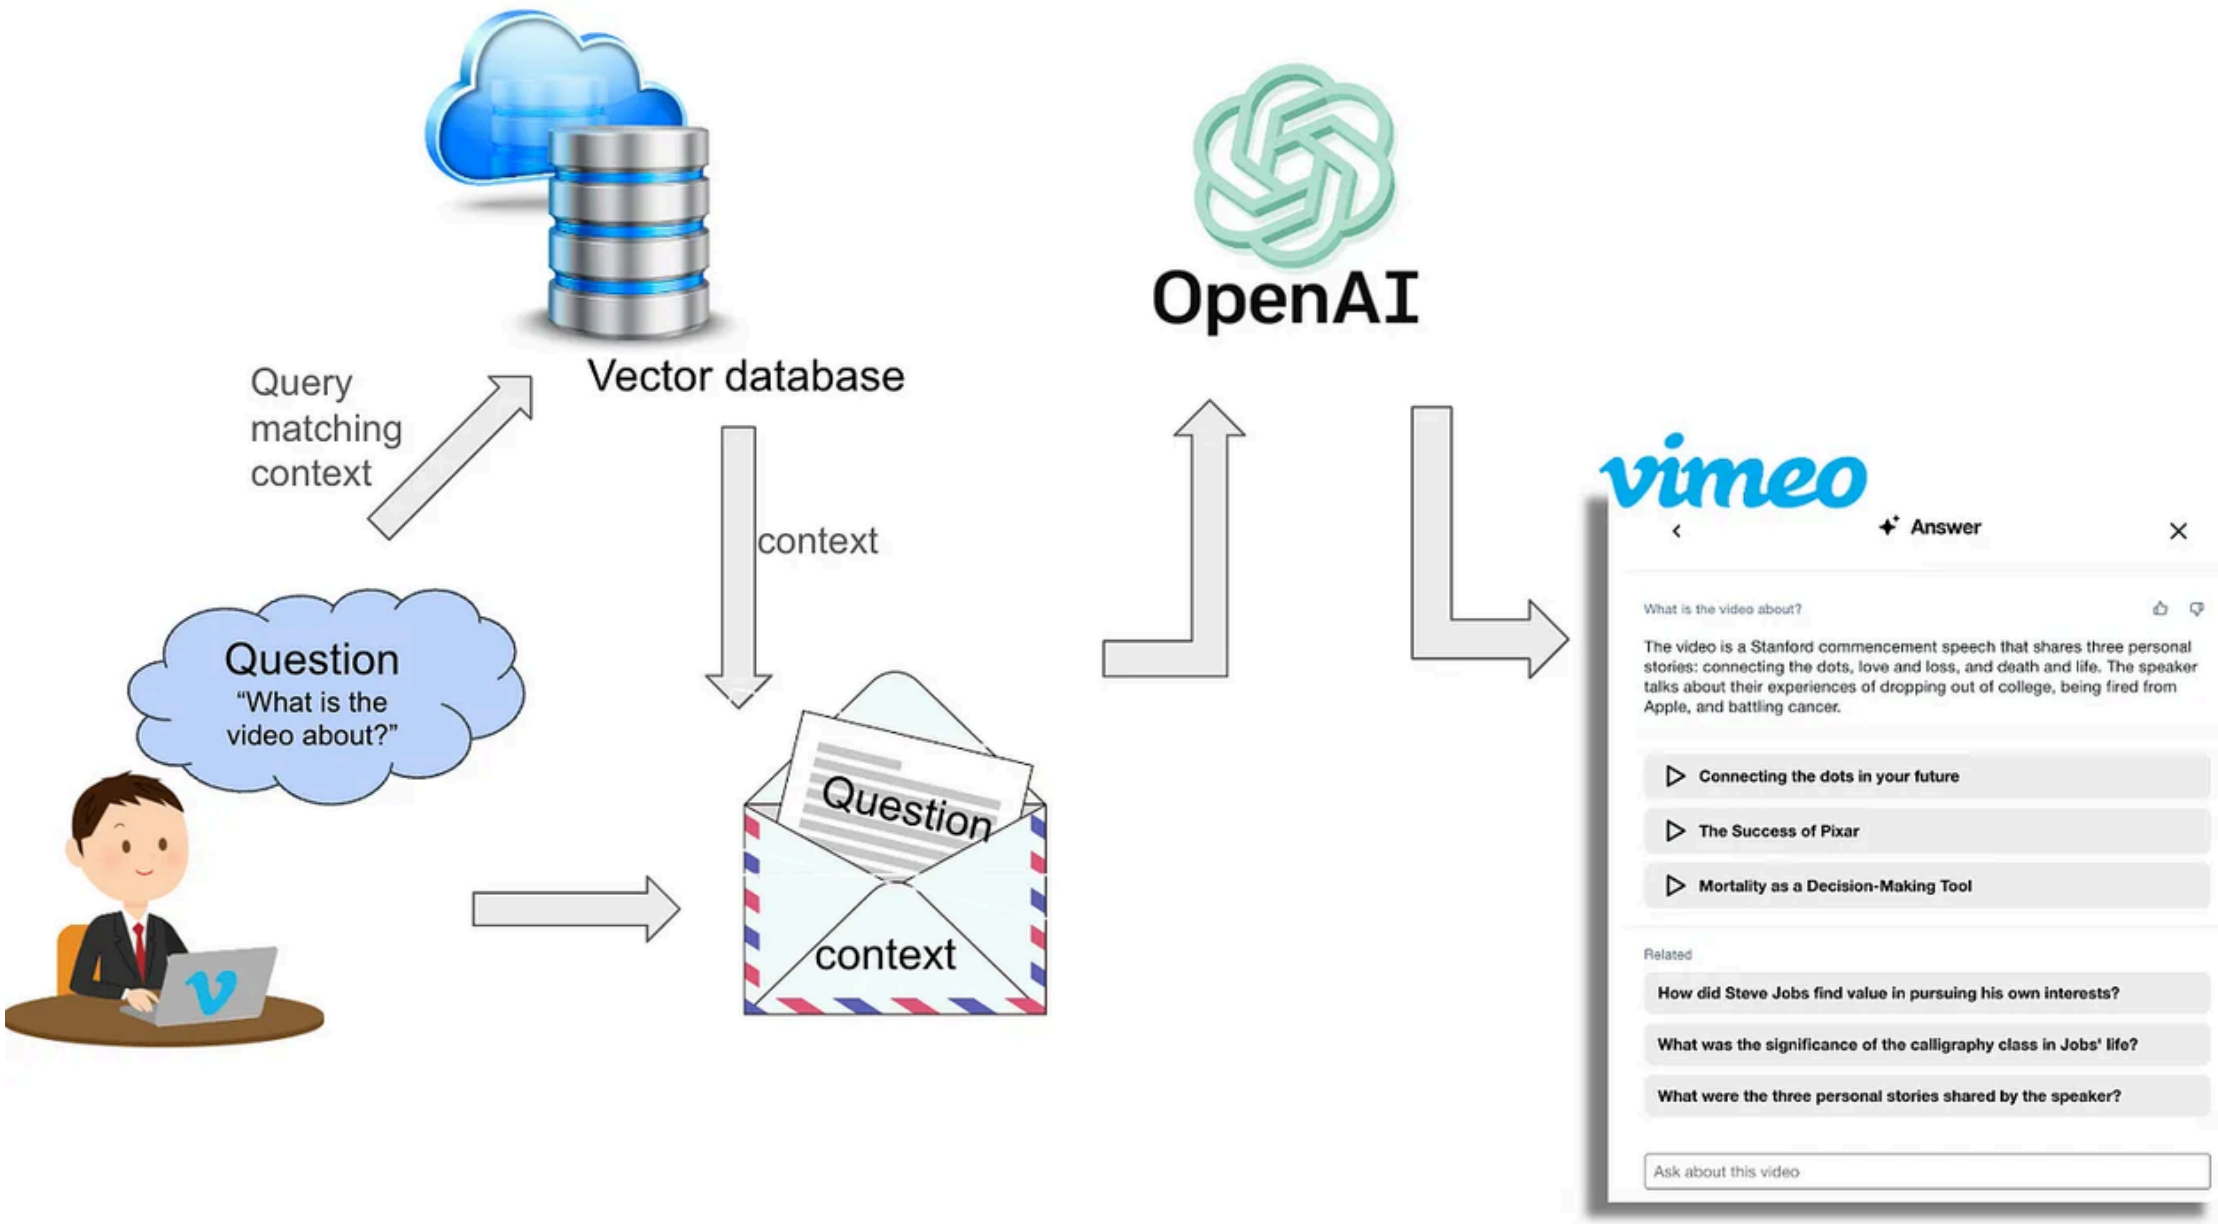

Vimeo built a RAG-based chatbot that lets users interact with videos in a conversational way. The chatbot can summarize videos, highlight key moments, and suggest follow-up questions.

The process starts by converting the video into text using its transcript. This transcript is then processed and stored in a vector database. Vimeo follows a bottom-up approach when registering transcripts: it uses multiple context window sizes, creates summaries for long sections, and generates an overall description for the full video.
When a user asks a question, such as What is the video about?, the system retrieves the most relevant transcript chunks from the vector database. These are passed to an LLM, which generates an answer grounded in the retrieved context.

Along with the response, the chatbot provides clickable video timestamps that support the answer. To make the interaction more engaging, it also suggests pre-generated question-and-answer pairs covering key moments in the video, as well as related questions based on the user's query.